

Linear Regression

Using the normal equation and scikit-learn

Brais Galvan Sotelo (19563)

CS550 - NPU

Table of Contents

- Introduction
- Normal Equation Implementation of linear regression
 - Set-up
 - Data description
 - Plotting the data
 - Theta Values and predictions
 - Plotting the results
- Scikit-learn Implementation of linear regression
 - Model Creation and prediction
 - Theta value calculation
- Results
- Conclusion

Implementation - Linear Regression using Normal Equation

Step 1.1: General Set-up

```
[53] # Python ≥3.5 is required
import sys
assert sys.version_info >= (3, 5)

# Scikit-Learn ≥0.20 is required
import sklearn
assert sklearn.__version__ >= "0.20"

# Common imports
import numpy as np
import os

# to make this notebook's output stable across runs
np.random.seed(42)
```

First we are going to import some necessary modules to complete the project, and make sure that Python version and Scikit-learn version are appropriate.

We also set-up a seed to get consistent result over multiple runs of the code. (This would be important if we use random generated data)

Implementation - Linear Regression using Normal Equation

Step 1.2: Figure creation Set-up

```
# To plot pretty figures
%matplotlib inline
import matplotlib as mpl
import matplotlib.pyplot as plt
mpl.rc('axes', labelsiz=14)
mpl.rc('xtick', labelsiz=12)
mpl.rc('ytick', labelsiz=12)

# Where to save the figures
PROJECT_ROOT_DIR = "."
CHAPTER_ID = "training_linear_models"
IMAGES_PATH = os.path.join(PROJECT_ROOT_DIR, "images", CHAPTER_ID)
os.makedirs(IMAGES_PATH, exist_ok=True)

def save_fig(fig_id, tight_layout=True, fig_extension="png", resolution=300):
    path = os.path.join(IMAGES_PATH, fig_id + "." + fig_extension)
    print("Saving figure", fig_id)
    if tight_layout:
        plt.tight_layout()
    plt.savefig(path, format=fig_extension, dpi=resolution)
```

We define how we want our images/figures to look and then we define where we want to save the images.

Implementation - Linear Regression using Normal Equation

Step 1.3: Data Set-up

```
[55] import numpy as np
import pandas as pd

from google.colab import files
uploaded = files.upload()
import io
abalone = pd.read_csv(
    io.BytesIO(uploaded['abalone_train.csv']),
    names=["Length", "Diameter", "Height", "Whole weight", "Shucked weight",
           "Viscera weight", "Shell weight", "Age"])

X1 = abalone["Length"]
X2 = np.array(X1)
X = X2.reshape(-1, 1)

y1 = abalone["Height"]
y2 = np.array(y1)
y = y2.reshape(-1, 1)
```

Choose Files | abalone_train.csv

- **abalone_train.csv**(application/vnd.ms-excel) - 149229 bytes, last modified: 6/7/2021 - 100% done
Saving abalone_train.csv to abalone_train (1).csv

We will upload a file containing the data for this problem. The data is contained in the file `abalone_train.csv`. The `abalone_train` data has 8 columns of information "Length", "Diameter", "Height", "Whole weight", "Shucked weight", "Viscera weight", "Shell weight" and "Age". (A sample of the data is shown in the next slide)

The we define X and Y to be Length and the Height

Abalone_train DATA Sample

	A	B	C	D	E	F	G	H	I
1	0.435	0.335	0.11	0.334	0.1355	0.0775	0.0965	7	
2	0.585	0.45	0.125	0.874	0.3545	0.2075	0.225	6	
3	0.655	0.51	0.16	1.092	0.396	0.2825	0.37	14	
4	0.545	0.425	0.125	0.768	0.294	0.1495	0.26	16	
5	0.545	0.42	0.13	0.879	0.374	0.1695	0.23	13	
6	0.57	0.45	0.145	0.751	0.2825	0.2195	0.2215	10	
7	0.47	0.36	0.13	0.472	0.182	0.114	0.15	10	
8	0.61	0.45	0.19	1.0805	0.517	0.2495	0.2935	10	
9	0.52	0.425	0.125	0.79	0.372	0.205	0.19	8	
10	0.485	0.39	0.12	0.599	0.251	0.1345	0.169	8	
11	0.625	0.495	0.155	1.025	0.46	0.1945	0.34	9	
12	0.615	0.495	0.16	1.255	0.5815	0.3195	0.3225	12	
13	0.455	0.35	0.14	0.5185	0.221	0.1265	0.135	10	
14	0.475	0.355	0.115	0.5195	0.279	0.088	0.1325	7	
15	0.385	0.3	0.1	0.2895	0.1215	0.063	0.09	7	
16	0.67	0.525	0.165	1.6085	0.682	0.3145	0.4005	11	


The abalone_train data has 3320 rows of information with 8 parameter each.

Implementation - Linear Regression using Normal Equation

Step 2: Plotting X and Y values

```
[58] plt.plot(X, y, "b.")  
     plt.xlabel("$x_1$", fontsize=18)  
     plt.ylabel("$y$", rotation=0, fontsize=18)  
     plt.axis([0, 1, 0, 1.5])  
     save_fig("generated_data_plot")  
     plt.show()
```

Saving figure generated_data_plot



Generate a figure by plotting the X and Y values. We define the axis in such a way that we accommodate the data properly, and then save the figure.

Implementation - Linear Regression using Normal Equation

Step 3: Finding the theta values

```
[60] X_b = np.c_[np.ones((3320, 1)), X] # add x0 = 1 to each instance
      theta_best = np.linalg.inv(X_b.T.dot(X_b)).dot(X_b.T).dot(y)
```

```
[61] theta_best
```

```
array([[ -0.0108267 ],
       [  0.28716253]])
```

```
[62] X_new = np.array([[0], [2]])
      X_new_b = np.c_[np.ones((2, 1)), X_new] # add x0 = 1 to each instance
      y_predict = X_new_b.dot(theta_best)
      y_predict
```

```
array([[ -0.0108267 ],
       [  0.56349837]])
```

We apply the normal equation to get theta values. In this case we use `np.linalg.inv()` in order to calculate the multiplicative inverse of the matrix.


$$\text{Theta} = (X^T X)^{-1} X^T Y$$

Then we use theta to make predictions of Y values using new values of X.

Implementation - Linear Regression using Normal Equation

Step 4.1: Plotting X and Y and the new data

```
[64] plt.plot(X_new, y_predict, "r-")  
     plt.plot(X, y, "b.")  
     plt.axis([0, 1, 0, 1.5])  
     plt.show()
```




Now in red, we can see the linear regression line, defined by plotting X_{new} and y_{pred} .

Implementation - Linear Regression using Normal Equation

Step 4.2: Adding additional features to the figure

```
[66] plt.plot(X_new, y_predict, "r-", linewidth=2, label="Predictions")
plt.plot(X, y, "b.")
plt.xlabel("$x_1$", fontsize=18)
plt.ylabel("$y$", rotation=0, fontsize=18)
plt.legend(loc="upper left", fontsize=14)
plt.axis([0, 1, 0, 1.5])
save_fig("linear_model_predictions_plot")
plt.show()
```

Saving figure linear_model_predictions_plot



Here we are simply going to add some additional features to our graph, including labels, for the axis and the legend for the regression prediction line (in red).

Implementation - Linear Regression using scikit-learn

Step 1: Creating a model and making a prediction

```
[68] from sklearn.linear_model import LinearRegression

lin_reg = LinearRegression()
lin_reg.fit(X, y)
lin_reg.intercept_, lin_reg.coef_

(array([-0.0108267]), array([[0.28716253]]))
```

```
[69] lin_reg.predict(X_new)

array([[ -0.0108267 ],
       [ 0.56349837]])
```

We utilize our previous definition of X and Y values together with the Linear Regression functionality from scikit-learn to get the intercept and coefficient of the regression equation, which we can then use to make a prediction to new values of X

Implementation - Linear Regression using scikit-learn

Step 2: Finding the theta values

The `LinearRegression` class is based on the `scipy.linalg.lstsq()` function (the name stands for "least squares"), which you could call directly:

```
[70] theta_best_svd, residuals, rank, s = np.linalg.lstsq(X_b, y, rcond=1e-6)
     theta_best_svd

array([[ -0.0108267 ],
       [  0.28716253]])
```

This function computes $\mathbf{X}^+ \mathbf{y}$, where \mathbf{X}^+ is the *pseudoinverse* of \mathbf{X} (specifically the Moore-Penrose inverse). You can use `np.linalg.pinv()` to compute the pseudoinverse directly:

```
np.linalg.pinv(X_b).dot(y)
```

```
array([[ -0.0108267 ],
       [  0.28716253]])
```

+ Code

+ Text

We use predefined functions to directly find the theta values,

Conclusion

In this scenario, we implemented the Normal Equation and scikit-learn to get a Linear Regression algorithm and make a prediction. Both method yielded the same answer.

Scikit-learn, provides the tools to get implemented the Linear Regression faster and more easily, but it can be somewhat abstract and hard to understand.

On the other hand, using the Normal Equation was more tedious but I gave me a better understanding of the results.

Also, plotting the data was extremely useful to both understand the data and double-check the result of the regression model among other things.